DECISION MEMO

Omnichannel Retail Performance Command Centre

To: Retail Executive Committee From: Harkirat Singh Date: 18 July 2026

Decision: Reallocate retention, store-diagnostic, and marketing attention using the validated business-case model.

Basis: CC BY 4.0 UCI Online Retail transactions combined with deterministic simulated stores, costs, sessions, campaigns, and attribution. Store and marketing results are simulated business-case evidence, not observed company performance.

1. Protect concentrated customer value; separate retention from reactivation

Champions are 714 of 4,371 identified customers (16.3%) but generate 46.1% of monetary value and average £1,793 observed historical CLV. At Risk and Lost segments contain 1,335 customers (30.5%). In the unsimulated UCI core, 65.6% of 4,338 identified purchasing customers place at least two completed orders.

Action: Protect Champions with service and loyalty treatment. Run a capped reactivation test for At Risk customers; suppress Lost customers after two failed contacts.

2. Diagnose matched-store gaps before changing network-wide targets

In the simulated large-store tier, Manchester is £94,542 (+18.0%) above its peer baseline while Birmingham is £94,542 (-18.0%) below. Leeds is £44,976 (+11.4%) above the standard-store baseline; Bristol is £27,233 (-6.9%) below.

Action: Compare each matched pair on category mix, returns, conversion, discounting, and staffing. Transfer only practices that survive those checks.

3. Stop simulated Social spend from absorbing margin

Simulated Social produces -47.0% ROI and £421 acquisition cost. Email produces +52.4% ROI at £168 per conversion; Paid Search produces +51.1% at £165.

Action: Reduce Social prospecting in the next simulated planning cycle, move the test budget to Email and Paid Search, and preserve a Social holdout.

Limitations

- The public source is historical and online-only; it does not represent current retail conditions.
- Physical stores, costs, categories, sessions, campaigns, spend, and attribution are simulated.
- Last non-direct attribution assigns credit; it does not estimate causal lift.
- Customer IDs are missing for 135,080 raw rows, so customer metrics exclude anonymous transactions.
- December 2011 contains nine days and is excluded from complete-month comparisons.
- DAX-to-SQL reconciliation covers six scenarios; other filter combinations remain untested.

Metrics to monitor

Decision area	Metrics	Cadence	Trigger
Customer value	Retention, repeat rate, segment migration, CLV, contribution/customer	Monthly	Two declines
Store operations	Revenue, margin rate, returns, peer variance by format/category	Weekly/monthly	Two negative periods
Marketing	Spend, conversion, CAC, attributed revenue, ROAS, ROI, holdout lift	Weekly/close	ROI < 0 or CAC > contribution

Evidence: CLM-002 to CLM-008 in 03_synthesis/evidence-register.md. Computed tables are under 02_analysis/outputs/tables/.

Source: Chen, D. (2015). Online Retail [Dataset]. UCI Machine Learning Repository. https://doi.org/10.24432/C5BW33. CC BY 4.0.